

IMPERIAL


agent-spec-kit

Scenario-Driven Evaluation and Testing of Tool-Using LLM Agents

Tool-using agents act on the world

- LLMs are no longer just text generators. As **agents** they plan, call tools, and **change external state**.
- Correctness can therefore **not** be judged from the final reply alone.
- A support agent can sound perfectly polite while it:
 - skips **authentication** and leaks account data,
 - calls the **wrong tool** or in the wrong order,
 - leaves the **database** in an invalid state.

The plausible answer is only one part of being correct.



The reply passes, but the actions and state silently do not.

Correctness lives on three surfaces

Output is necessary but far from sufficient



Output

what the agent says



Trace

which tools it called



State

the state it left behind

A failure on any surface can be hidden by a fluent reply, and may even be masked by recovery on a later turn.

Existing tools only cover part of the loop

- **Output graders / LLM-judges**

- Good at final-answer grading
- Blind to tool & state failures

- **Agent benchmarks**


- Fixed leaderboards
- Not reusable infrastructure for custom agents

- **Trace viewers**

- Show what happened
- Not what should have happened

The gap

No single, reusable way to assert over **output + trace + state** in one multi-turn scenario, with **precise diagnostics** when a check fails.



Trace viewers show the episode, not the expected behaviour.


...and you can't script every conversation

Authored tests only cover what you thought to write

- Hand-written scenarios capture the policies you already know: the cases you thought of.
- But agents are **stochastic** and real users are **messy**: they rephrase, omit details, push back, change intent mid-conversation.
- The failures that bite are the dialogue paths **nobody scripted**.

Beyond the cases you authored

The framework should **explore** conversations automatically (simulated users and transcript fuzzing), and **keep** the real failures it finds as permanent regression tests.



Scripts cover a small known region. Bugs hide in the rest.

Research question

Can the evaluation of **stateful, tool-using LLM agents** be expressed in a **software-testing style**:

readable scenarios that combine assertions over **replies, tool trajectories**, and **application state**, while keeping enough structured evidence to **diagnose** failures precisely?

Answer: agent-spec-kit, a scenario-driven testing framework for LLM agents, evaluated on a 50-task benchmark.

Part I: The Library

agent-spec-kit: behaviour as executable scenarios

The core idea: a scenario is a conversation script

Output, tool, and state checks live together, woven in after each turn

```
@scenario(agent_fixture="agent", user_fixture="user")
async def test_billing_credit(s, store):
    s.user_message("I was double charged in October")
    # O: the reply must insist on verifying first
    s.assert_output(m.contains("verify"))
    # T: the agent must authenticate before account tools
    s.assert_tool_calls([
        m.tool_call("authenticate_customer"),
        m.tool_call("get_billing", args={"id": "CUST-045"}),
    ])
    # S: a postcondition on the database state
    s.assert_that(assert_credit_policy)
```

- **O** assert_output: the reply
- **T** assert_tool_calls / forbid_tool_calls: the trace
- **S** assert_that: environment / DB state

One file reads like a dialogue with the contract inlined, with no separate datasets, graders, or runner config.

One composable matcher language

Start simple: matching the reply

- Literals, dicts and lists are coerced into matchers automatically.
- Combine them: `m.one_of` (any), `m.all_of` (all), `m.not_`, `m.contains`, `m.regex`.
- Selective **LLM rubric** only where text needs judgement: pass/fail criteria, not opaque scores.

```
# deterministic: offer a refund OR a credit
s.assert_output(
    m.one_of(
        m.contains("refund"),
        m.contains("account credit"),
    )
)
```

```
# nuance? hand that one check to an LLM judge
s.assert_output(m.llm_criteria(
    criteria=["acknowledges the double charge",
              "does not promise a refund yet"],
    threshold=2,
))
```

Matching the tool trace

Order and extras are explicit, not implied

- `assert_tool_calls` matches the normalised tool trace.
- `ordered=True` pins the sequence, `ordered=False` allows parallel siblings.
- `allow_extras=True` lets other calls interleave.
- `forbid_tool_calls` asserts a tool is never used.

```
s.assert_tool_calls(  
  [  
    m.tool_call("authenticate_customer"),  
    m.tool_call("get_billing",  
                 args={"id": "CUST-045"}),  
  ],  
  ordered=True, # must appear in order  
  allow_extras=True, # others may interleave  
)
```

...and the matchers compose

Structured objects and conditional rules

- `m.object` constrains structured args or results, matching only the fields you name.
- `m.list` matches sequences ordered or unordered.
- **Conditional rules:**
require/forbid a field only when another field holds.

```
m.tool_call("open_incident", args=m.object(  
    {  
        "id": m.regex(r"INC-\d{4}-\d+"),  
        "severity": m.one_of("high", "medium"),  
        "components": m.list(["edge", "core"],  
                             mode="unordered"),  
    },  
    rules=[ # conditional  
        m.require("pager_group")  
        .when(m.field("severity") == "high"),  
        m.forbid("pager_group")  
        .when(m.field("severity") == "medium"),  
    ],  
))
```

Precise diagnostics: structured counterexamples

When a check fails, you learn which expectation broke, where

- Every failed assertion produces a **Counterexample**:
 - headline + human location (“assert_tool_calls after turn #2”)
 - a **JSON-pointer path** to the mismatch, e.g. `$(1).args.line_id`
 - expected vs. actual, plus a bounded event trace
- Identical view in the **terminal**, the **SQLite** store, and the **web UI**.
- Turns “it failed somehow” into “this check broke here”.

Typical grader output

✗ AssertionError / score 0.0

which check? which turn? which field?

agent-spec-kit counterexample


[0] authenticate ✓

[1] get_line_status ✗

path `$(1).args.line_id`
expected LINE-001
actual LINE-002


A scalar tells you that it failed. A counterexample tells you where.

Framework-neutral execution model



- Adapters **normalise traces** (LangGraph, Pydantic AI) into one event model, so oracles never depend on a framework's message format.

From authored tests to automatic discovery



- **Repeats** for stochastic agents.
- **User simulation:** persona-driven multi-turn dialogues.
- **Mutation fuzzing:** perturb valid transcripts.

- **Shrinking:** delta-debug to a minimal failing dialogue.
- **Extraction:** pin a discovered failure as a permanent scenario.
- All under the same oracle vocabulary.

Persistent runs, CLI, and a web results browser

- Every run stored locally (SQLite + gzip blobs) with git metadata.
- CLI: run, runs, show, regressions, ui.
- Read-only **web UI**: runs list, run detail, trace drawer, fuzz trials.
- **Compare** named experiments side-by-side to track regressions.

TRIALS	SCENARIO	TAGS	REPEAT	STATUS	DURATION	FAILURE KIND	FAILURE MESSAGE	ASSERTIONS
—	test_007_c29_full	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.F MULTIANTFAULT_WRONG_CLARIFICATION	Q/1	failed	10.79s	forbid_tool_calls	forbid_tool_calls: test_007_c29_full: forbidtool tool call 'order_replacement_loan' found at index 0 (partial lo...	Q/1
—	test_007_c29_output	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.D MULTIANTFAULT_WRONG_CLARIFICATION	1/1	passed	61.09s	—	—	Q/0
—	test_007_c29_state	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.S MULTIANTFAULT_WRONG_CLARIFICATION	1/1	passed	33.97s	—	—	Q/0
—	test_007_c29_trace	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.T MULTIANTFAULT_WRONG_CLARIFICATION	Q/1	failed	10.12s	forbid_tool_calls	forbid_tool_calls: test_007_c29_trace: forbidtool tool call 'order_replacement_loan' found at index 0 (partial lo...	Q/1
—	test_005_340_full	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.F MULTIANTFAULT_WRONG_NESTED_TOOL	Q/1	failed	1m 14s	assert_tool_calls	assert_tool_calls: test_005_340_full: could not match expected element at index 1 while preserving order...	Q/1
—	test_005_340_output	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.D MULTIANTFAULT_WRONG_NESTED_TOOL	1/1	passed	1m 37s	—	—	Q/0
—	test_005_340_state	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.S MULTIANTFAULT_WRONG_NESTED_TOOL	1/1	passed	1m 2s	—	—	Q/0
—	test_005_340_trace	TELECOM FAULT-DETECTION FAULT7FEE TASK7F3 ORACLE.T	Q/1	failed	1m 7s	assert_tool_calls	assert_tool_calls: test_005_340_trace: could not match expected element at index 1 while preserving...	Q/1

Run summary failed
SCENARIOS 65 / 152 passed
REPEATS 65 / 152 passed
REPEAT PASS RATE 42.8%
FLAKY SCENARIOS 0
MIN REPORT 56.13s
MAX REPORT 53.88s / 1m 58s
STARTED 26/05/2026, 16:29:13
FINISHED 26/05/2026, 16:48:15

Run metadata
Run ID run_20260526_152913_38e403
Experiment default
Python 3.14.0
Package 61.0
Git fault-seeding-calibration @ 7948c5 (d87c)
Command agent-spec-kit run /Users/ruthe/.../agent_spec_kit/benchmark.../testcase_support/tasks/Fault_.../detect.py --tags Fault -detection -v 8

Failure kinds
assert_tool_calls x24
assert_output x13
assert_error x10
forbid_tool_calls x5

Tag breakdown

Run detail page: per-scenario, per-repeat results with a trace drawer.

Live Demo

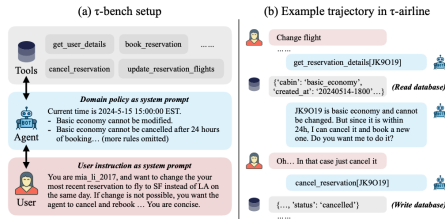
Writing a scenario, watching it fail precisely,
discovering a bug, and pinning it as a regression.

Part II: Evaluation

Does multi-surface, scenario-driven testing actually pay off?

TelcoSupportBench: the evaluation testbed


- A synthetic **mobile-support** domain built for this project.
- **50 tasks** (T01–T50) against a reference **LangGraph** agent: a coordinator + diagnostics + billing specialists.
- Each task has **four oracle lanes**: **O** output, **T** trace, **S** state, and **F** = full (all three).
- Compared against **7** other evaluation approaches.



Inspiration: τ -bench-style dual-control support setting. [\[Replace with TelcoSupportBench architecture figure if available\]](#)

Finding 1: output-only grading is not enough

Reference agent across 200 oracle slots (50 tasks $\times$ 4 lanes)




- **All 50 output lanes passed**, yet **54 of 200** slots failed.
- Every one of those failures lived in the **trace**, **state**, or **full** lanes.
- The reply “sounded correct” while a tool ran in the wrong order, an argument was wrong, or a store update never happened.

Takeaway

Final-answer evaluation alone would have reported **100%** pass.

Finding 2: more expressive, less code

Twelve canonical check patterns, ported across frameworks



- **165 lines** vs. 496 (DeepEval) and 532 (Pydantic Evals).
- **Highest clarity grade on every one** of the twelve specimens.
- On failures, it gives an exact **path witness** where alternatives return a scalar, broad prose, or a bare `AssertionError`.

Declarative DSL, equivalent in policy to the hand-written ports, but far easier to read and debug.

Finding 3: full oracles catch faults narrow lanes miss

Injected-fault detection (F01--F10) on previously-passing slots

- Ten fault families injected into the agent, then measure which lane detects each.
- **Output** detection is frequently **0%** for tool-ordering and state-mutation faults.
- **State** and **Full** lanes reach **75–100%** on those same faults.
- The **combined oracle** is stronger than any single surface alone.

Fault	O	S	T	F
F02 (state)	0%	100%	67%	100%
F03 (state)	0%	100%	0%	100%
F05 (order)	0%	0%	100%	100%
F07 (trace)	50%	50%	100%	100%
F08 (trace)	0%	0%	75%	50%
F10 (full)	25%	0%	33%	100%

Detection rate = detected / eligible. Selected rows. Full matrix F01–F10 in the report.

Finding 4: generative testing finds bugs nobody scripted

14 tasks, the same F-oracle throughout

Method	Convs	Tool paths	Failure sigs
Manual scripts	14	10	3
User simulation	70	58	24
Mutation fuzz	70	36	18

- Simulation & fuzzing reach **far more failure signatures** than the hand-written scripts, at little extra authoring cost.
- Discovered failures are **shrunk** and **extracted** into permanent regressions.

Threats to validity

- **Construct:** clarity & diagnostic grades used **LLM judges**, not a human developer study.
- **External:** a **single** reference agent and **one** telecom domain, so design choices may favour the framework.
- **Rerun:** extraction produced a valid file for every eligible failure, but **fewer than half** reproduced the exact signature on immediate rerun, as stochastic ordering and wording drift.
- Results are **promising** but describe this setup, not agents in general.

Conclusion

- The right unit of evaluation is the **behavioural episode**, not the final reply: a polite answer can hide a skipped auth step, a mis-ordered tool, or a corrupted store.
- agent-spec-kit unifies **output + trace + state** oracles in one readable, executable, framework-neutral scenario.
- **Diagnostics are part of evaluation quality**: assertion-local counterexamples make failures fixable, not just countable.
- A single workflow spans **discovery** → **regression**: simulate, fuzz, shrink, extract.
- Evidence: hidden trace/state failures exposed, **3× less** check code with the highest clarity, and generative testing surfacing real bugs, such as a pre-auth privacy leak.

Future work

- More **domains and agent frameworks**, to test whether the scenario model transfers beyond TelcoSupportBench.
- **Human studies** of authoring clarity and diagnostic usefulness, to replace LLM-judge grading.
- **Signature-stable replay** after extraction, and more robust matching of stochastic trajectories.
- **Larger, team-based suites** to see how the workflow scales in everyday development.

IMPERIAL

Thank you.
Questions?